

Deep Learning: Day 3

chyld @ galvanize

Topics

— — —

- Install Python Modules
- Download Training Datasets
- Setup Kaggle
- Setup AWS EC2 for Deep Learning
- Image Classification with Real World data
- Transfer Learning
- Image Classification using Transfer Learning
- Multi-label Image Classification

Packages to Install

- `pip install -U scikit-learn pillow glances`
- `pip install -U opencv-python tqdm seaborn kaggle`

Kaggle: Dogs vs Cats

<https://www.kaggle.com/c/dogs-vs-cats>

Get train.zip



Dogs vs. Cats

Create an algorithm to distinguish dogs from cats

215 teams · 4 years ago

[Overview](#) [Data](#) [Discussion](#) [Leaderboard](#) [Rules](#) [Team](#)

Competition Data

 sampleSubmission.csv

 test1.zip

 train.zip

train.zip 543.16 MB


 Download

Kaggle: Understanding the Amazon from Space

Featured Prediction Competition

Planet: Understanding the Amazon from Space

Use satellite data to track the human footprint in the Amazon rainforest

 Planet · 938 teams · a year ago

\$60,000

Prize Money

Overview

Data

Kernels

Discussion

Leaderboard


Rules


Team


My Submissions


Late Submission


Competition Data


 Kaggle-planet-test-t...


 Kaggle-planet-train-...


 sample_submission_v2...


 test-jpg-additional...

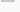
 test-jpg.tar.7z

 test-tif-v2.tar.7z

 test_v2_file_mapping...


 train-jpg.tar.7z

 train-tif-v2.tar.7z

 train_v2.csv.zip

test_v2_file_mapping.csv.zip

110.24 KB

 Download

https://www.kaggle.com/c/planet-understanding-the-amazon-from-space

Get train_v2.csv.zip & train-jpg.tar.7z

Kaggle

— — —

- `https://www.kaggle.com`
 - Register for account
- Get API key
 - `https://www.kaggle.com/<username>/account`
 - `mv kaggle.json ~/.kaggle`
- Install Kaggle CLI
 - `pip install -U kaggle`

Download Dogs vs Cats Images from Kaggle

— — —

- `https://www.kaggle.com/c/dogs-vs-cats`
- Accept rules
- `https://www.kaggle.com/c/dogs-vs-cats/rules`
- Download images using CLI/Terminal
- `"kaggle competitions download -c dogs-vs-cats"`

Amazon Web Services



Deep Learning AMI (Ubuntu) Version 10.0 - ami-e580c79d

Free tier eligible

Comes with latest binaries of deep learning frameworks pre-installed in separate virtual environments: MXNet, TensorFlow, Caffe, Caffe2, PyTorch, Keras, Chainer, Theano and CNTK. Fully-configured with NVidia CUDA, cuDNN and NCCL as well as Intel MKL-DNN

Root device type: ebs Virtualization type: hvm ENA Enabled: Yes

Select

64-bit

Model	NVIDIA Tesla V100 GPUs	GPU Memory	NVIDIA NVLink	vCPUs	Main Memory	Network Bandwidth	EBS Bandwidth
p3.2xlarge	1	16 GiB	n/a	8	61 GiB	Up to 10 Gbps	1.5 Gbps
p3.8xlarge	4	64 GiB	200 GBps	32	244 GiB	10 Gbps	7 Gbps
p3.16xlarge	8	128 GiB	300 GBps	64	488 GiB	25 Gbps	14 Gbps

NVIDIA Tesla V100 GPU



HP nVidia Tesla V100 16GB PCIE x16 876340-001 876908-001 Q2N68A

[Be the first to review this product...](#) [Ask Or Answer A Question](#)

In stock. Ships from United States. Most customers

Sold and Shipped by [Compeve](#)

- Manufacturer nVidia
- Model Tesla V100
- MPN 876340-001
- SKU 101000100503
- Warranty 30 days
- UPC 190017178622
- Content Video card, power cord
- Condition New, pulls



Sold and Shipped by: **Compeve**

\$8,899.⁰⁰

\$20.00 Shipping (restrictions apply)

1

+

-

ADD TO CART ►



ADD TO COMPARE



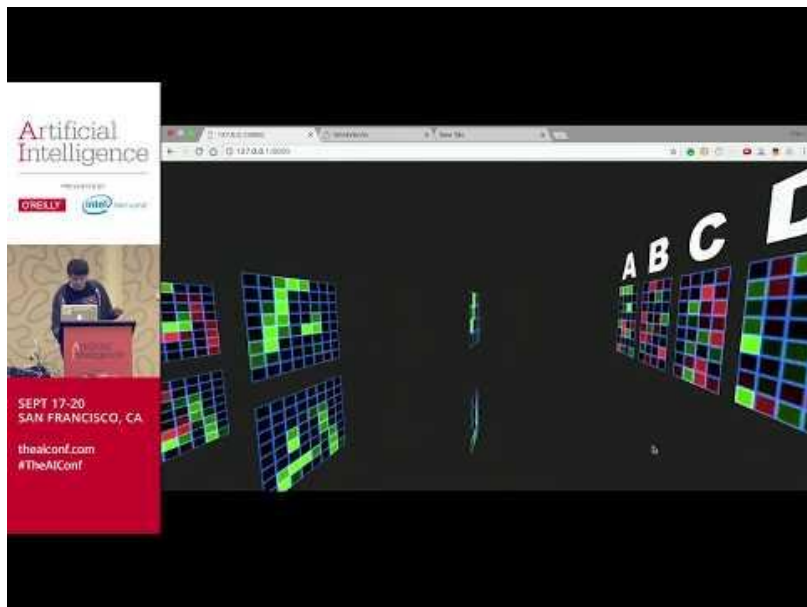
PRICE ALERT





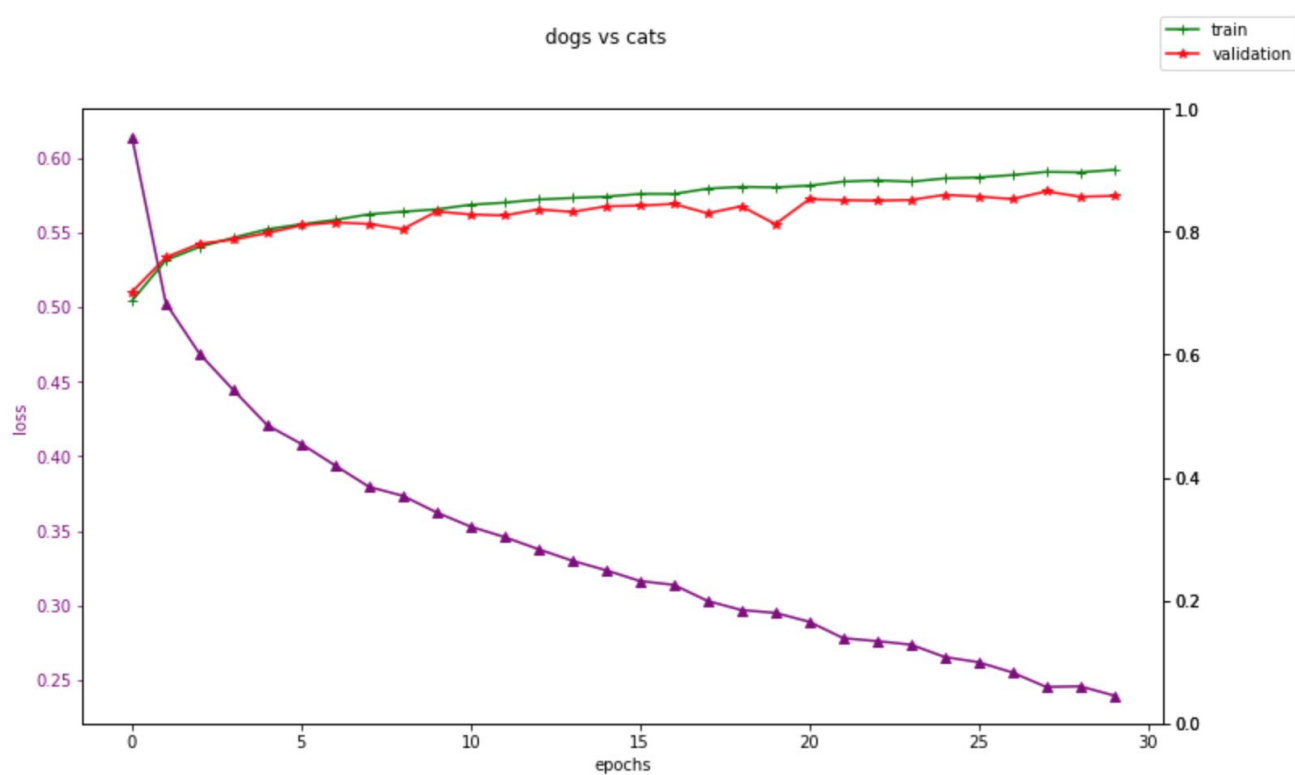
ADD TO WISH LIST

Review: Visualizing ConvNet Architecture

- https://youtu.be/0qm9vsf_hvU?t=265



CNN: Real World Data



Transfer Learning

— — —

What is Transfer Learning?

Transfer learning, is a research problem in machine learning that focuses on storing knowledge gained while solving one problem and applying it to a different but related problem.


Transfer Learning

Why Transfer Learning?


- In practice a very few people train a Convolution network from scratch (random initialisation) because it is rare to get enough dataset. So, using pre-trained network weights as initialisations or a fixed feature extractor helps in solving most of the problems in hand.
- Very Deep Networks are expensive to train. The most complex models take weeks to train using hundreds of machines equipped with expensive GPUs.
- Determining the topology/flavour/training method/hyper parameters for deep learning is a black art with not much theory to guide you.

Transfer Learning

Transfer Learning with CNNs



1. Train on ImageNet



2. If small dataset: fix all weights (treat CNN as fixed feature extractor), retrain only the classifier

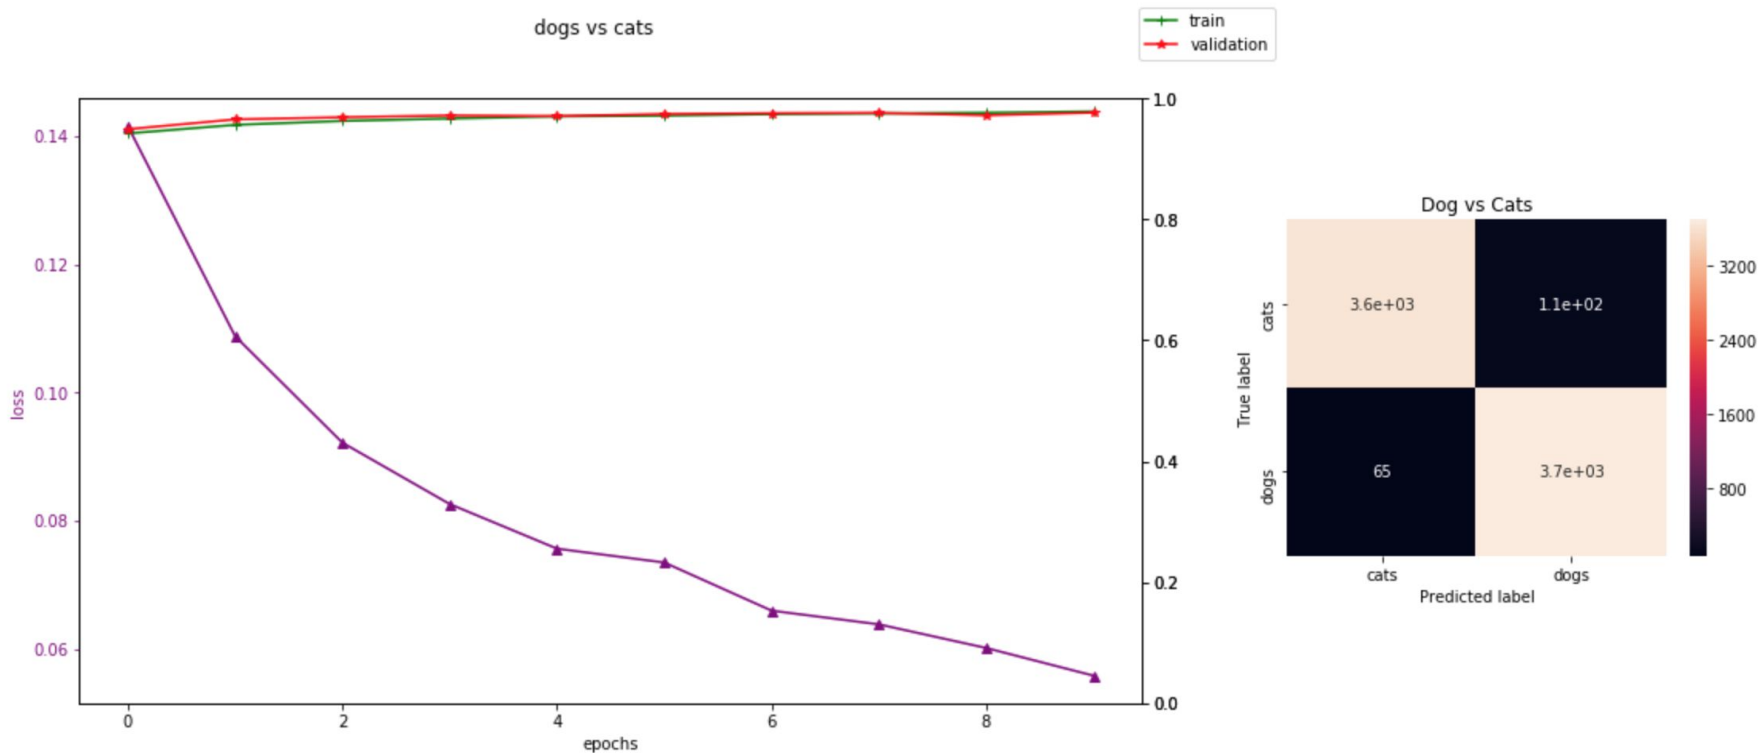
i.e. swap the Softmax layer at the end



3. If you have medium sized dataset, “**finetune**” instead: use the old weights as initialization, train the full network or only some of the higher layers


retrain bigger portion of the network, or even all of it.

Transfer Learning



Multi-Label Classification

```
[(0.999877, 'clear'), (0.9303916, 'agriculture'), (0.9999151, 'primary'), (0.6120964, 'road')]  
[(1, 'clear'), (1, 'agriculture'), (1, 'primary'), (1, 'road')]
```



```
[(1.0, 'clear'), (0.90402347, 'agriculture'), (1.0, 'primary'), (0.6447628, 'water')]  
[(1, 'clear'), (1, 'agriculture'), (1, 'primary'), (1, 'cultivation'), (1, 'water')]
```

